

# I2ML :: BASICS

## Data

$\mathcal{X} \subseteq \mathbb{R}^p$  :  $p$ -dimensional **feature space** / input space  
Usually we assume categorical features to be numerically encoded.

$\mathcal{Y}$  : **target space**  
e.g.:  $\mathcal{Y} = \mathbb{R}$  for regression,  $\mathcal{Y} = \{0, 1\}$  or  $\mathcal{Y} = \{-1, +1\}$  for binary classification,  $\mathcal{Y} = \{1, \dots, g\}$  for multi-class classification with  $g$  classes

$\mathbf{x} = (x_1, \dots, x_p)^T \in \mathcal{X}$  : **feature vector** / covariate vector

$y \in \mathcal{Y}$  : **target variable** / output variable  
Concrete samples are called labels

$(\mathbf{x}^{(i)}, y^{(i)}) \in \mathcal{X} \times \mathcal{Y}$  :  $i$ -th **observation** / sample / instance / example

$\mathbb{D} = \bigcup_{n \in \mathbb{N}} (\mathcal{X} \times \mathcal{Y})^n$  : **set of all finite data sets**

$\mathbb{D}_n = (\mathcal{X} \times \mathcal{Y})^n \subseteq \mathbb{D}$  : **set of all finite data sets of size  $n$**

$\mathcal{D} = ((\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(n)}, y^{(n)})) \in \mathbb{D}_n$  : **data set** of size  $n$ . An  $n$ -tuple, a family indexed by  $\{1, \dots, n\}$ . We use  $\mathcal{D}_n$  to emphasize its size.

$\mathcal{D}_{\text{train}}, \mathcal{D}_{\text{test}} \subseteq \mathcal{D}$  : **data sets for training and testing**  
Often:  $\mathcal{D} = \mathcal{D}_{\text{train}} \dot{\cup} \mathcal{D}_{\text{test}}$

$\mathbb{P}_{xy}$  : **joint probability distribution on  $\mathcal{X} \times \mathcal{Y}$**

### Classification

$o_k(y) = \mathbb{I}(y = k) \in \{0, 1\}$ : multiclass one-hot encoding, if  $y$  is class  $k$   
We write  $\mathbf{o}(y)$  for the  $g$ -length encoding vector and  $o_k^{(i)} = o_k(y^{(i)})$

$\pi_k = \mathbb{P}(y = k)$ : **prior probability** for class  $k$   
In case of binary labels we might abbreviate:  $\pi = \mathbb{P}(y = 1)$ .

## Model and Learner

**Model** / Hypothesis:  $f : \mathcal{X} \rightarrow \mathbb{R}^g$  maps features to predictions, often parametrized by  $\boldsymbol{\theta} \in \Theta$  (then we write  $f_{\boldsymbol{\theta}}(\mathbf{x})$  or  $f(\mathbf{x}|\boldsymbol{\theta})$ ).

$\Theta \subseteq \mathbb{R}^d$  : **parameter space**

$\boldsymbol{\theta} = (\theta_1, \theta_2, \dots, \theta_d) \in \Theta$ : model **parameter** vector  
Some models may traditionally use different symbols.

$\mathcal{H} = \{f : \mathcal{X} \rightarrow \mathbb{R}^g \mid f \text{ belongs to a certain functional family} \}$  :  
**Hypothesis space** – set of functions to which we restrict learning

**Learner** / Inducer  $\mathcal{I} : \mathbb{D} \times \boldsymbol{\Lambda} \rightarrow \mathcal{H}$  takes a training set  $\mathcal{D}_{\text{train}} \in \mathbb{D}$ , produces model  $f : \mathcal{X} \rightarrow \mathbb{R}^g$ , with hyperparam. configuration  $\boldsymbol{\lambda} \in \boldsymbol{\Lambda}$ .  
We also write  $\mathcal{I} : \mathbb{D} \times \boldsymbol{\Lambda} \rightarrow \Theta$  or  $\mathcal{I}_{\boldsymbol{\lambda}} : \mathbb{D} \rightarrow \Theta$

$\boldsymbol{\Lambda} = \boldsymbol{\Lambda}_1 \times \boldsymbol{\Lambda}_2 \times \dots \times \boldsymbol{\Lambda}_{\ell} \subseteq \mathbb{R}^{\ell}$ : **hyperparameter space**  
 $\boldsymbol{\Lambda}_j$  are usually bounded real or integer intervals or a finite categorical set

$\boldsymbol{\lambda} = (\lambda_1, \lambda_2, \dots, \lambda_{\ell}) \in \boldsymbol{\Lambda}$  : **hyperparameter configuration**

$\epsilon = y - f(\mathbf{x})$  or  $\epsilon^{(i)} = y^{(i)} - f(\mathbf{x}^{(i)})$  : ( $i$ -th) **residual** in regression

### Classification

$\pi_k(\mathbf{x}) : \mathcal{X} \rightarrow [0, 1]$  **probability prediction** for class  $k$ , approximates  $\mathbb{P}(y = k \mid \mathbf{x})$ ; for binary we abbreviate with  $\pi(\mathbf{x})$  for  $\mathbb{P}(y = 1 \mid \mathbf{x})$ .

$f_k(\mathbf{x}) : \mathcal{X} \rightarrow \mathbb{R}$ : **scoring** / discriminant **function** for class  $k$ ;  
for binary we use  $f(\mathbf{x}) = f_1(\mathbf{x}) - f_2(\mathbf{x})$

$h(\mathbf{x}) : \mathcal{X} \rightarrow \mathcal{Y}$  : **hard label function**;  
Typically created by  $h(\mathbf{x}) = \arg \max_{k \in \{1, \dots, g\}} f_k(\mathbf{x})$  or  $h(\mathbf{x}) = \arg \max_{k \in \{1, \dots, g\}} \pi_k(\mathbf{x})$

$y f(\mathbf{x})$  or  $y^{(i)} f(\mathbf{x}^{(i)})$ : **margin** for ( $i$ -th) observation in binary classification

$\hat{y}, \hat{f}, \hat{h}, \hat{\pi}_k(\mathbf{x}), \hat{\pi}(\mathbf{x})$  and  $\hat{\boldsymbol{\theta}}$   
The hat symbol denotes **learned** functions and parameters.

## Loss, Risk and ERM

$L : \mathcal{Y} \times \mathbb{R}^g \rightarrow \mathbb{R}_0^+$  : **loss function**: Quantifies "quality"  $L(y, f(\mathbf{x}))$  of prediction  $f(\mathbf{x})$  (or  $L(y, \pi(\mathbf{x}))$  of prediction  $\pi(\mathbf{x})$ ) for true  $y$ .

$\mathcal{R} : \mathcal{H} \rightarrow \mathbb{R}$  : **(theoretical) risk** ;  $\mathcal{R}(f) = \mathbb{E}_{((\mathbf{x}, y) \sim \mathbb{P}_{xy})}[L(y, f(\mathbf{x}))]$

$\mathcal{R}_{\text{emp}} : \mathcal{H} \rightarrow \mathbb{R}$  : **empirical risk** ;  $\mathcal{R}_{\text{emp}}(f) = \sum_{i=1}^n L(y^{(i)}, f(\mathbf{x}^{(i)}))$

Since  $f$  is usually defined by **parameters  $\boldsymbol{\theta}$** , we also write:  
 $\mathcal{R}_{\text{emp}} : \Theta \rightarrow \mathbb{R}$ ;  $\mathcal{R}_{\text{emp}}(\boldsymbol{\theta}) = \sum_{i=1}^n L(y^{(i)}, f(\mathbf{x}^{(i)} \mid \boldsymbol{\theta}))$

**Empirical risk minimization (ERM)**:  $\hat{\boldsymbol{\theta}} \in \arg \min_{\boldsymbol{\theta} \in \Theta} \mathcal{R}_{\text{emp}}(\boldsymbol{\theta})$

## Regression Losses

- L2 loss / squared error:**
- $L(y, f(\mathbf{x})) = (y - f(\mathbf{x}))^2$  or  $L(y, f(\mathbf{x})) = 0.5(y - f(\mathbf{x}))^2$
  - Convex and differentiable, non-robust against outliers
  - Optimal constant model:  $\hat{f}(\mathbf{x}) = \frac{1}{n} \sum_{i=1}^n y^{(i)} = \bar{y}$
  - Optimal model over  $\mathbb{P}_{xy}$  for unrestricted  $\mathcal{H}$ :  $\hat{f}(\mathbf{x}) = \mathbb{E}[y|\mathbf{x}]$

- L1 loss / absolute error:**
- $L(y, f(\mathbf{x})) = |y - f(\mathbf{x})|$
  - Convex and more robust, non-differentiable
  - Optimal constant model:  $\hat{f}(\mathbf{x}) = \text{med}(y^{(1)}, \dots, y^{(n)})$
  - Optimal model over  $\mathbb{P}_{xy}$  for unrestricted  $\mathcal{H}$ :  $\hat{f}(\mathbf{x}) = \text{med}[y|\mathbf{x}]$

## Classification Losses

**0-1-loss (binary case)**  
 $L(y, h(\mathbf{x})) = \mathbb{I}(y \neq h(\mathbf{x}))$   
 $L(y, f(\mathbf{x})) = \mathbb{I}(y f(\mathbf{x}) < 0)$  for  $\mathcal{Y} = \{-1, +1\}$   
Discontinuous, results in NP-hard optimization

**Brier score (binary case)**  
 $L(y, \pi(\mathbf{x})) = (\pi(\mathbf{x}) - y)^2$  for  $\mathcal{Y} = \{0, 1\}$   
Least-squares on probabilities

**Log-loss / Bernoulli loss / binomial loss (binary case)**  
 $L(y, \pi(\mathbf{x})) = -y \log(\pi(\mathbf{x})) - (1 - y) \log(1 - \pi(\mathbf{x}))$  for  $\mathcal{Y} = \{0, 1\}$   
 $L(y, \pi(\mathbf{x})) = \log(1 + (\frac{\pi(\mathbf{x})}{1 - \pi(\mathbf{x})})^{-y})$  for  $\mathcal{Y} = \{-1, +1\}$

Assuming a logit-link  $\pi(\mathbf{x}) = \exp(f(\mathbf{x})) / (1 + \exp(f(\mathbf{x})))$ :  
 $L(y, f(\mathbf{x})) = -y \cdot f(\mathbf{x}) + \log(1 + \exp(f(\mathbf{x})))$  for  $\mathcal{Y} = \{0, 1\}$   
 $L(y, f(\mathbf{x})) = \log(1 + \exp(-y \cdot f(\mathbf{x})))$  for  $\mathcal{Y} = \{-1, +1\}$   
Penalizes confidently-wrong predictions heavily

**Brier score (multi-class case)**

$L(y, \pi(\mathbf{x})) = \sum_{k=1}^g (\pi_k(\mathbf{x}) - o_k(y))^2$

**Log-loss (multi-class case)**

$L(y, \pi(\mathbf{x})) = - \sum_{k=1}^g o_k(y) \log(\pi_k(\mathbf{x}))$

Optimal constant models  
0-1-loss:  $h(\mathbf{x}) \in \arg \max_{j \in \{0, 1\}} \sum_{i=1}^n \mathbb{I}(y^{(i)} = j)$

Brier and log-loss (binary):  $\hat{\pi}(\mathbf{x}) = \bar{y}$

Brier and log-loss (multiclass):  $\hat{\pi}(\mathbf{x}) = \left( \frac{1}{n} \sum_{i=1}^n o_1^{(i)}, \dots, \frac{1}{n} \sum_{i=1}^n o_g^{(i)} \right)$