Exercise 1: Car Price Prediction

Imagine you work at a second-hand car dealer and are tasked with finding for-sale vehicles your company can acquire at a reasonable price. You decide to address this challenge in a data-driven manner and develop a model that predicts adequate market prices (in EUR) from vehicles' properties.

a) Characterize the task at hand: supervised or unsupervised? Regression or classification? Learning to explain or learning to predict? Justify your answers.

b) How would you set up your data? Name potential features along with their respective data type and state the target variable.

c) Assume now that you have data on vehicles' age (days), mileage (km), and price (EUR). Explicitly define the feature space $\mathcal{X}$ and target space $\mathcal{Y}$.

d) You choose to use a linear model (LM) for this task. For this, you assume the targets to be conditionally independent given the features, i.e., $y^{(i)}|\mathbf{x}^{(i)} \perp y^{(j)}|\mathbf{x}^{(j)}$ for all $i, j \in \{1, 2, \ldots, n\}, i \neq j$, with sample size n. The LM models the target as a linear function of the features with Gaussian error term: $\mathbf{y} = \mathbf{X}\boldsymbol{\theta} + \epsilon$,
$\epsilon \sim N(\mathbf{0}, diag(\sigma^2)), \quad \sigma > 0$.

State the hypothesis space for the corresponding model class. For this, assume the parameter vector $\boldsymbol{\theta}$ to include the intercept coefficient.

e) Which parameters need to be learned? Define the corresponding parameter space Θ.

f) State the loss function for the i-th observation using $L2$ loss.

g) In classical statistics, you would estimate the parameters via maximum likelihood estimation (MLE). The likelihood for the LM is given by:

$$\mathcal{L}(\boldsymbol{\theta}|\mathbf{x}) = \prod_{i=1}^{n} \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{1}{2\sigma^2} \left(y^{(i)} - \boldsymbol{\theta}^T\mathbf{x}^{(i)} \right)^2 \right)$$

Describe how you can make use of the likelihood in empirical risk minimization (ERM) and write down the resulting empirical risk.

h) Now you need to optimize this risk to find the best parameters, and hence the best model, via empirical risk minimization. State the optimization problem formally and list the necessary steps to solve it.

Congratulations, you just designed your first machine learning project!