


Exercise 1: HRO in mlr3

Throughout the lecture, we will frequently use the R package `mlr3` and its descendants, providing an integrated ecosystem for all common machine learning tasks. Let's recap the HRO principle and see how it is reflected in `mlr3`. An overview of the most important objects and their usage, illustrated with numerous examples, can be found at <https://mlr3book.ml-org.com/basics.html>.


- a) How are the key concepts you learned about in the lecture videos implemented in `mlr3`?
- b) Have a look at `mlr3::tsk("iris")`. What attributes does this `task` object store?
- c) Pick an `mlr3` learner of your choice. What are the different settings for this learner?
(Hint: Use `mlr3::mlr_learners$keys()` to see all available learners.)

Exercise 2: Loss Functions for Regression Tasks

In this exercise, we will examine loss functions for regression tasks somewhat more in-depth.



- a) Consider the above linear regression task. How will the model parameters be affected by adding the new outlier point (orange) if you use
 - i) $L1$ loss
 - ii) $L2$ lossin your empirical risk? (You do not need to actually compute the parameter values.)



b) The second plot visualizes another loss function popular in regression tasks, the so-called *Huber loss* (depending on $\delta > 0$; here: $\delta = 5$). Describe how the Huber loss deals with residuals as compared to $L1$ and $L2$ loss. Can you guess its definition?

c) **Zu schwer?** Show that $median(y)$ is the optimal constant prediction c when using $L1$ loss.
Hint:

– Employing the law of total expectation, we can find c via

$$\arg \min_c \mathbb{E}[|y - c|] = \arg \min_c \int_{-\infty}^{\infty} |y - c| p(y) dy = \arg \min_c \int_{-\infty}^c -(y - c) p(y) dy + \int_c^{\infty} (y - c) p(y) dy.$$

– Setting the derivative of this expression, found by application of Leibniz's rule, to zero yields


$$0 \stackrel{!}{=} \int_{-\infty}^c p(y) dy - \int_c^{\infty} p(y) dy.$$

d) Derive the least-squares estimator, i.e., the solution to the linear model when using $L2$ loss, analytically via

$$\arg \min_{\theta} \|y - \mathbf{X}\theta\|_2^2.$$

Exercise 3: Polynomial Regression

Assume the following (noisy) data-generating process: $y = 0.5 + \sin(2\pi x) + \epsilon$ with $\epsilon \sim \mathcal{N}(0, 0.1)$.



- a) We decide to model the data with a cubic polynomial (including intercept term). State the corresponding hypothesis space.
- b) Demonstrate that this hypothesis space is simply a parameterized family of curves by plotting in R curves for 3 different models belonging to the considered model class.
- c) State the empirical risk w.r.t. θ for a member of our hypothesis space. Use $L2$ loss and be as explicit as possible.
- d) We can minimize this risk using gradient descent. In order to make this somewhat easier, we will denote the transformed feature matrix, containing x to the power from 0 to 3, by $\tilde{\mathbf{X}}$, such that we can express our model by $\tilde{\mathbf{X}}\theta$ (note that the model is still linear in its parameters, even if \mathbf{X} has been transformed in a non-linear manner!). Derive the gradient of the empirical risk w.r.t θ .
- e) Using the result from d), state the calculation to update the current parameter $\theta^{[t]}$.
- f) You will not be able to fit the data perfectly with a cubic polynomial. Should you opt for a more flexible model class (i.e., a hypothesis space with higher capacity)? What might be disadvantageous about this?

Exercise 4: Predicting abalone

We want to predict the age of an abalone using its longest shell measurement and its weight.

See <http://archive.ics.uci.edu/ml/datasets/Abalone> for more details.

```
url <-
  "http://archive.ics.uci.edu/ml/machine-learning-databases/abalone/abalone.data"
abalone <- read.table(url, sep = ",", row.names = NULL)
colnames(abalone) <- c(
  "sex", "longest_shell", "diameter", "height", "whole_weight",
  "shucked_weight", "visceral_weight", "shell_weight", "rings")
```

- a) Plot `longest_shell` and `whole_weight` on the x - and y -axis, respectively, and color points with `rings`.

Using `mlr3`:

- b) Define a linear regression learner (for this you will need to load the `mlr3learners` extension package first) and use it to fit a linear model to the `abalone` data. For this, first create a `task` like so (necessary if they have heard the intro to `mlr3` by then?):

```
data <- abalone[, c("longest_shell", "shell_weight", "rings")]
task_abalone <- mlr3::TaskRegr$new(
  id = "abalone",
  backend = data,
  target = "rings")
```

- c) Compare the fitted and observed targets visually.
(Hint: use `autoplot()`.)
- d) Assess the model's training loss in terms of MAE.
(Hint: losses are retrieved by calling `$score()`, which accepts different `mlr_measures`, on the prediction object.)

Recall that you can find the official `mlr3` manual at <https://mlr3book.ml-org.com/>.