Solution 1:

In the following, you assume that your outcome follows a $\log_2$-normal distribution with density function

$$p(y|f) = \frac{1}{y\sigma\sqrt{2\pi}} \exp\left(-\frac{(\log_2(y) - f)^2}{2\sigma^2}\right).$$

In other words, $\log_2(Y)$ follows a normal distribution. You observe $n = 3$ data points $\boldsymbol{y}$ and want to model f using features $\boldsymbol{X} \in \mathbb{R}^{n \times p}$. You choose to use a gradient boosting tree algorithm.

(a)
- $\ell(f) = const - (\log_2(y) - f)^2/2\sigma^2$ (1P)
- $\partial\ell/\partial f = \sigma^{-2}(\log_2(y) - f) \propto (\log_2(y) - f)$ (1P)

(b) Use $\tilde{y} = \log_2(y) = (0, 1, 2)$ (1P)

(i) $\hat{f}^{[0]}(\boldsymbol{x}) = \bar{\tilde{y}} = 1$ as this is the optimal constant model for squared error. (1P)

(ii) $\boldsymbol{r}^{[1]} = \sigma^{-2}(-1, 0, 1)$ (1P)

(iii) $R_t^{[1]}, t = 1, 2$ will split using $\boldsymbol{x}_1$, as $\boldsymbol{x}_2$ carries no information. Since $x_1^{(1)} = x_1^{(2)}$,

$$R_1 = -0.5I(x_1 \geq 0.5)$$

and

$$R_2 = 1I(x_1 \leq 0.5).$$

(2P)

(iv) $\hat{f}^{[1]}(\boldsymbol{x}) = (0.5, 0.5, 2)$ (1P)

(v) $\boldsymbol{r}^{[2]} = \sigma^{-2}(-0.5, 0.5, 0)$ (1P)

(c) Nothing, because there is no information that can be used to further improve the model. (1P)

(d) (i) M grows: The capacity will increase and the algorithm may eventually overfit (1P)

(ii) n grows: The capacity will stay the same and the algorithm may underfit (1P)