

Solution 1:

a) Benchmark result:

(i) Total number of models trained:

$$\underbrace{4 \cdot 10}_{\text{outer resampling}} + 2 \cdot \underbrace{10 \cdot \underbrace{5 \cdot 200}_{\text{one tuning iteration}}}_{\substack{\text{all outer folds in one tuning procedure} \\ \text{all tuning procedures}}} = 20,040.$$

(ii) Since we evaluate on AUC, we select k -NN with the best average result in that respect.

b) Less data for training leads to higher bias, less data for evaluation leads to higher variance.

c) Statements:

- i) True – 3-CV leads to smaller train sets, therefore we are not able to learn as well as in, e.g., 10-CV.
- ii) False – we are relatively flexible in choosing the outer loss, but the inner loss needs to be suitable for empirical risk minimization, which encompasses differentiability in most cases (i.e., whenever optimization employs derivatives).

Solution 2:

This exercise is a compact version of a tutorial on `mlr3gallery`. Feel free to explore the additional steps and explanations featured in the original (there is also a bunch of other useful code demos).

```
a) library(mlr3verse)
library(mlr3tuning)

## Loading required package: mlr3
## Loading required package: paradox

(task <- tsk("pima"))

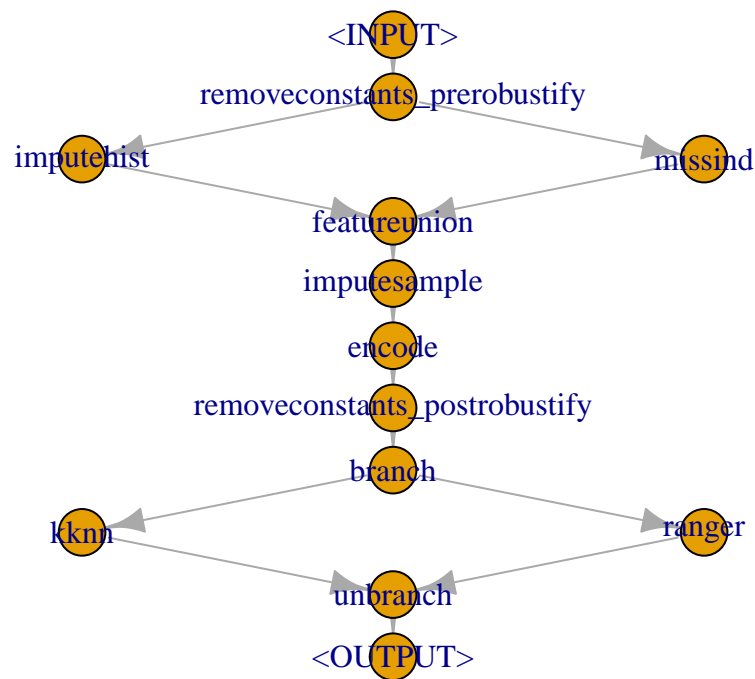
## <TaskClassif:pima> (768 x 9)
## * Target: diabetes
## * Properties: twoclass
## * Features (8):
##   - dbl (8): age, glucose, insulin, mass, pedigree, pregnant, pressure,
##     triceps
```

```
b) learners <- list(
  po(lrn("classif.kknn", id = "kknn")),
  po(lrn("classif.ranger", id = "ranger")))
```

```
c) ppl_preproc <- ppl("robustify", task = task, factors_to_numeric = TRUE)
```

```
d) ppl_learners <- ppl("branch", learners)
```

```
e) ppl_combined <- ppl_preproc %>% ppl_learners  
plot(ppl_combined)
```



```
graph_learner <- as_learner(ppl_combined)
```

```
f) # check available hyperparameters for tuning (converting to data.table for
# better readability)
tail(as.data.table(graph_learner$param_set), 10)
```

##		id	class	lower	upper	levels	nlevels
## 1:		ranger.oob.error	ParamLgl	NA	NA	TRUE,FALSE	2
## 2:		ranger.max.depth	ParamInt	-Inf	Inf		Inf
## 3:		ranger.alpha	ParamDbl	-Inf	Inf		Inf
## 4:		ranger.min.prop	ParamDbl	-Inf	Inf		Inf
## 5:	ranger.regularization.factor	ParamUty		NA	NA		Inf
## 6:	ranger.regularization.usedepth	ParamLgl		NA	NA	TRUE,FALSE	2
## 7:		ranger.seed	ParamInt	-Inf	Inf		Inf
## 8:		ranger.minprop	ParamDbl	-Inf	Inf		Inf
## 9:		ranger.se.method	ParamFct	NA	NA	jack,infjack	2
## 10:		branch.selection	ParamInt	1	2		2
##	is_bounded	special_vals	default	storage_type		tags	
## 1:	TRUE	<list[0]>	TRUE	logical		train	
## 2:	FALSE	<list[1]>		integer		train	
## 3:	FALSE	<list[0]>	0.5	numeric		train	
## 4:	FALSE	<list[0]>	0.1	numeric		train	
## 5:	FALSE	<list[0]>	1	list		train	
## 6:	TRUE	<list[0]>	FALSE	logical		train	
## 7:	FALSE	<list[1]>		integer		train	
## 8:	FALSE	<list[0]>	0.1	numeric		train	
## 9:	TRUE	<list[0]>	infjack	character		predict	
## 10:	TRUE	<list[0]>	<NoDefault[3]>	integer	train,predict,required		

```

# seeing all our hyperparameters of interest are of type int, we specify the
# tuning objects accordingly, and dependencies for k and mtry
graph_learner$param_set$values$branch.selection <-
  to_tune(p_int(1, 2))
graph_learner$param_set$values$kknn.k <-
  to_tune(p_int(3, 10, depends = branch.selection == 1))
graph_learner$param_set$values$ranger.mtry <-
  to_tune(p_int(1, 5, depends = branch.selection == 2))

# rename learner (otherwise, mlr3 will display a lengthy chain of operations
# in result tables)
graph_learner$id <- "graph_learner"

```

```
g) # make sure to set a seed for reproducible results
set.seed(123)

# perform nested resampling, terminating after 3 evaluations
rr <- tune_nested(
  method = "random_search",
  task = task,
  learner = graph_learner,
  inner_resampling = rsmp("cv", folds = 3),
  outer_resampling = rsmp("cv", folds = 3),
  measure = msr("classif.ce"),
  term_evals = 3)

```

```
h) rr$score()

##           task task_id      learner      learner_id
## 1: <TaskClassif[49]>   pima <AutoTuner[41]> graph_learner.tuned
## 2: <TaskClassif[49]>   pima <AutoTuner[41]> graph_learner.tuned
## 3: <TaskClassif[49]>   pima <AutoTuner[41]> graph_learner.tuned
##           resampling resampling_id iteration      prediction
## 1: <ResamplingCV[19]>         cv           1 <PredictionClassif[20]>
## 2: <ResamplingCV[19]>         cv           2 <PredictionClassif[20]>
## 3: <ResamplingCV[19]>         cv           3 <PredictionClassif[20]>
##   classif.ce
## 1:  0.2500000
## 2:  0.2421875
## 3:  0.2148438

rr$aggregate()

## classif.ce
##  0.2356771
```

The performance estimate for our tuned learner then amounts to an MCE of around 0.24.

Solution 3:

We do not provide an explicit solution here, but have a look at the [tuning code demo](#), which covers some parts, and take inspiration from the public contributions on Kaggle.