


# Introduction to Machine Learning

## Gradient Boosting - Illustration



### Learning goals

- Understand impact of different loss functions and
- Understand impact of different base learners for regression

# GRADIENT BOOSTING ILLUSTRATION - GAM

We now compare different loss functions and base learners. We start with a GAM as base learner and compare the  $L2$  loss with the  $L1$  loss.


**Reminder:** Pseudo-residuals

- $L2$ :  $\tilde{r}(f) = r(f) = y - f(\mathbf{x})$
- $L1$ :  $\tilde{r}(f) = \text{sign}(y - f(\mathbf{x}))$

We consider a regression task with a single feature  $x$  and target  $y$ , with the following true underlying relationship:

$$y^{(i)} = -1 + 0.2 \cdot x^{(i)} + 0.1 \cdot \sin(x^{(i)}) + \epsilon^{(i)}$$

$$\text{with } n = 50 \text{ and } \epsilon^{(i)} \sim \mathcal{N}(0, 0.1) \quad \forall i \in \{1, \dots, n\}$$

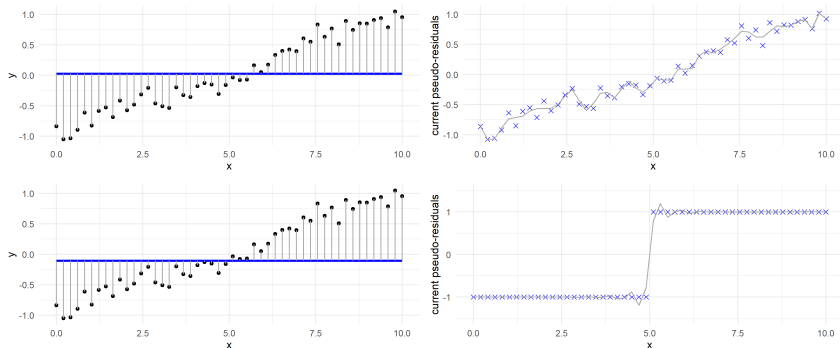


# GRADIENT BOOSTING ILLUSTRATION - GAM

- ❶ We start with the simplest model, the optimal constant – mean of the target variable in the case of  $L2$  loss and median in the case of  $L1$  loss.
- ❷ We improve the model by calculating the pointwise pseudo-residuals on the training data, and fit a GAM on the residuals.
  - ❶ The GAM base learners model the conditional mean via cubic  $B$ -splines with 40 knots.
  - ❷ In each step, the GAM fitted on the current pseudo-residuals is multiplied by a constant learning rate of 0.2 and added to the previous model.
  - ❸ This procedure is repeated multiple times.

# GAM WITH $L_2$ VS $L_1$ LOSS

Top:  $L_2$  loss, bottom:  $L_1$  loss

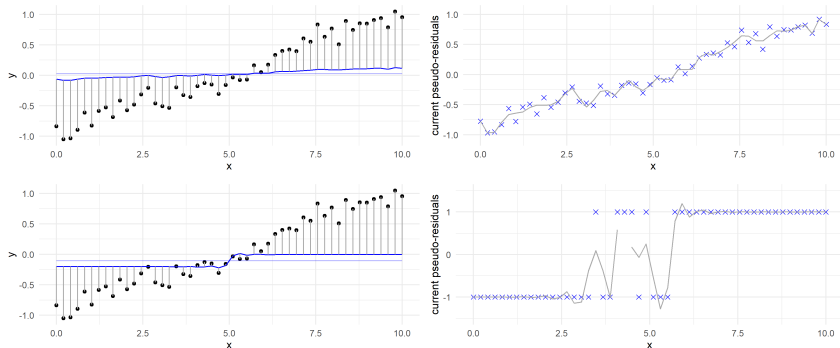


## Iteration 1

The nature of the pseudo-residuals affects gradual model fit: as  $L_1$  only considers residuals' sign, the corresponding base learners are less affected by very large or small values compared to  $L_2$  and hence lead to more moderate changes.

# GAM WITH $L_2$ VS $L_1$ LOSS

Top:  $L_2$  loss, bottom:  $L_1$  loss




## Iteration 2

The nature of the pseudo-residuals affects gradual model fit: as  $L_1$  only considers residuals' sign, the corresponding base learners are less affected by very large or small values compared to  $L_2$  and hence lead to more moderate changes.

# GAM WITH $L_2$ VS $L_1$ LOSS

Top:  $L_2$  loss, bottom:  $L_1$  loss




## Iteration 3

The nature of the pseudo-residuals affects gradual model fit: as  $L_1$  only considers residuals' sign, the corresponding base learners are less affected by very large or small values compared to  $L_2$  and hence lead to more moderate changes.

# GAM WITH $L_2$ VS $L_1$ LOSS

Top:  $L_2$  loss, bottom:  $L_1$  loss




Iteration 10

The nature of the pseudo-residuals affects gradual model fit: as  $L_1$  only considers residuals' sign, the corresponding base learners are less affected by very large or small values compared to  $L_2$  and hence lead to more moderate changes.

# GAM WITH $L_2$ VS $L_1$ LOSS

Top:  $L_2$  loss, bottom:  $L_1$  loss



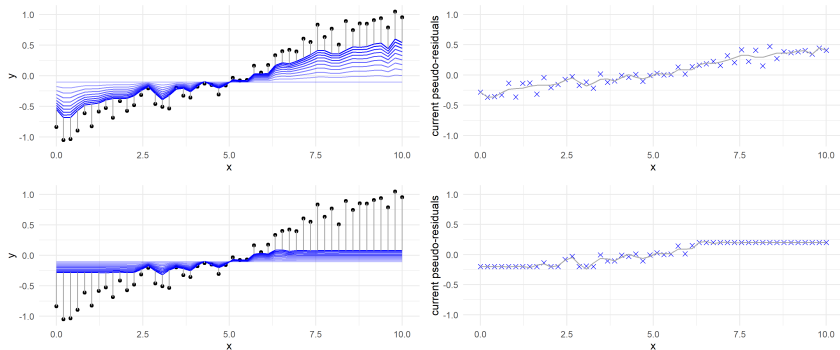
Iteration 100

The nature of the pseudo-residuals affects gradual model fit: as  $L_1$  only considers residuals' sign, the corresponding base learners are less affected by very large or small values compared to  $L_2$  and hence lead to more moderate changes.




# GAM WITH HUBER LOSS

We can also use Huber loss, which is closer to  $L_2$  for large  $\delta$  values and closer to  $L_1$  for smaller  $\delta$  values. Top:  $\delta = 2$ , bottom:  $\delta = 0.2$ .




## Iteration 10

We see how for smaller  $\delta$  (bottom) pseudo-residuals are often effectively bounded, resulting in  $L_1$ -like behavior, while the upper plot more closely resembles  $L_2$  loss.



# GAM WITH OUTLIERS

Instead of normally distributed noise we can assume a  $t$ -distribution, leading to outliers in the observed target values. Top:  $L_2$ , bottom:  $L_1$ .




## Iteration 10

$L_2$  loss is affected by outliers rather strongly, whereas  $L_1$  solely considers residuals' sign and not their magnitude, resulting in a more robust model.

# GAM WITH OUTLIERS

Instead of normally distributed noise we can assume a  $t$ -distribution, leading to outliers in the observed target values. Top:  $L_2$ , bottom:  $L_1$ .




Iteration 100

$L_2$  loss is affected by outliers rather strongly, whereas  $L_1$  solely considers residuals' sign and not their magnitude, resulting in a more robust model.

# LM WITH $L_2$ VS $L_1$ LOSS

Instead of using a GAM as base learner we now use a simple linear model. Top:  $L_2$ , bottom:  $L_1$ .

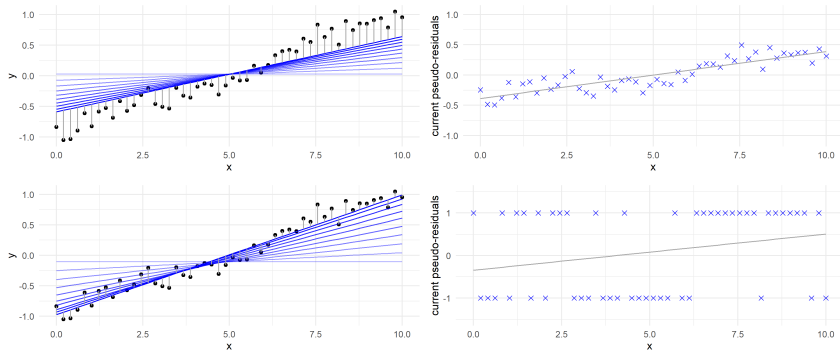


## Iteration 1

For  $L_2$ , as  $\tilde{r}(f) = r(f)$ , we find the optimal model in the very first iteration; only the multiplicative learning rate slows down optimization. In the  $L_1$  case the base learner LMs fit pseudo-residuals that differ from model residuals, leading to a less monotonic optimization path.

# LM WITH $L_2$ VS $L_1$ LOSS

Instead of using a GAM as base learner we now use a simple linear model. Top:  $L_2$ , bottom:  $L_1$ .




## Iteration 10

For  $L_2$ , as  $\tilde{r}(f) = r(f)$ , we find the optimal model in the very first iteration; only the multiplicative learning rate slows down optimization. In the  $L_1$  case the base learner LMs fit pseudo-residuals that differ from model residuals, leading to a less monotonic optimization path.

# LM WITH $L_2$ VS $L_1$ LOSS

Instead of using a GAM as base learner we now use a simple linear model. Top:  $L_2$ , bottom:  $L_1$ .



## Iteration 100

For  $L_2$ , as  $\tilde{r}(f) = r(f)$ , we find the optimal model in the very first iteration; only the multiplicative learning rate slows down optimization. In the  $L_1$  case the base learner LMs fit pseudo-residuals that differ from model residuals, leading to a less monotonic optimization path.

# LM: GB VS GD

- As we have seen, boosting with LMs and  $L_2$  loss simply approaches the closed-form OLS solution with a speed determined by the learning rate  $\beta$ .
- This is perfectly equivalent to fitting an LM via **gradient descent**, as a gradient step in parameter space is a gradient step in function space for this specific case.
- Recall the parameter update for GD with learning rate  $\beta$ :  
$$\boldsymbol{\theta}^{[m+1]} \leftarrow \boldsymbol{\theta}^{[m]} - \beta \cdot \nabla_{\boldsymbol{\theta}^{[m]}} \mathcal{R}_{\text{emp}}(\boldsymbol{\theta}^{[m]}) = \boldsymbol{\theta}^{[m]} + \beta(-\mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \boldsymbol{\theta}^{[m]}).$$
- Now compute the update for a boosted LM with current model  $\mathbf{X}\boldsymbol{\theta}^{[m]}$  (note that adding a linear base learner to an LM is equivalent to summing parameters):

$$\frac{\partial}{\partial \boldsymbol{\theta}^{[m+1]}} \left\| (\mathbf{y} - \mathbf{X}\boldsymbol{\theta}^{[m]}) - \mathbf{X}\boldsymbol{\theta}^{[m+1]} \right\|_2^2 = 0$$

$$-2\mathbf{X}^T (\mathbf{y} - \mathbf{X}\boldsymbol{\theta}^{[m]}) + 2\mathbf{X}^T \mathbf{X} \boldsymbol{\theta}^{[m+1]} = 0$$

$$\boldsymbol{\theta}^{[m+1]} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T (\mathbf{y} - \mathbf{X}\boldsymbol{\theta}^{[m]})$$

$$\boldsymbol{\theta}^{[m+1]} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} - (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{X} \boldsymbol{\theta}^{[m]}$$

$$\boldsymbol{\theta}^{[m+1]} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} - \boldsymbol{\theta}^{[m]}$$

$$\boldsymbol{\theta}^{[m+1]} = -\mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \boldsymbol{\theta}^{[m]}.$$

$$\Rightarrow \hat{f}^{[m+1]} = \mathbf{X}\tilde{\boldsymbol{\theta}}^{[m+1]} = \mathbf{X} \left( \boldsymbol{\theta}^{[m]} + \beta(-\mathbf{X}^T \mathbf{y} + \mathbf{X}^T \mathbf{X} \boldsymbol{\theta}^{[m]}) \right).$$