


Introduction to Machine Learning

Pseudo-Residuals and Gradient Descent



Learning goals


- Know the concept of pseudo-residuals
- Understand the relationship between pseudo-residuals and gradient descent

PSEUDO-RESIDUALS

- In regression, residuals are defined as $r := y - f(\mathbf{x})$.
- We further define **pseudo-residuals** as the negative first derivatives of loss functions w.r.t. $f(\mathbf{x})$

$$\tilde{r} := -\frac{\partial L(y, f(\mathbf{x}))}{\partial f(\mathbf{x})}.$$

- Note that pseudo-residuals are actually functions of y and $f(\mathbf{x})$ and depend on the loss function L .




BEST POINT-WISE UPDATE

Assume we have fitted a model $f(\mathbf{x})$ to data \mathcal{D} .

Assume we could update $f(\mathbf{x})$ point-wise as we like. For a fixed $\mathbf{x} \in \mathcal{X}$, the best point-wise update is the direction of the residual $r = y - f(\mathbf{x})$

$$f(\mathbf{x}) \leftarrow f(\mathbf{x}) + r$$




BEST POINT-WISE UPDATE

Assume we have fitted a model $f(\mathbf{x})$ to data \mathcal{D} .

Assume we could update $f(\mathbf{x})$ point-wise as we like. For a fixed $\mathbf{x} \in \mathcal{X}$, the best point-wise update is the direction of the residual $r = y - f(\mathbf{x})$

$$f(\mathbf{x}) \leftarrow f(\mathbf{x}) + r$$

The point-wise error at this specific \mathbf{x} becomes 0.



APPROXIMATE BEST POINT-WISE UPDATE

When applying gradient descent to compute a point-wise update of $f(\mathbf{x})$, we would go a step into the direction of the negative gradient

$$f(\mathbf{x}) \leftarrow f(\mathbf{x}) - \frac{\partial L(y, f(\mathbf{x}))}{\partial f(\mathbf{x})}.$$

which is the direction of the pseudo-residual

$$f(\mathbf{x}) \leftarrow f(\mathbf{x}) + \tilde{r}$$

Iteratively stepping towards the direction of the pseudo-residuals is the underlying idea of gradient boosting, which is a learning algorithm that will be covered in a later chapter.

GD IN ML AND PSEUDO-RESIDUALS

- In gradient descent, we try to move in the direction of the negative gradient in each step by updating the model accordingly

$$\boldsymbol{\theta}^{[t+1]} = \boldsymbol{\theta}^{[t]} - \alpha^{[t]} \cdot \nabla_{\boldsymbol{\theta}} \mathcal{R}_{\text{emp}}(\boldsymbol{\theta})|_{\boldsymbol{\theta}=\boldsymbol{\theta}^{[t]}}$$

with step size $\alpha^{[t]}$.

- This can be seen as approximating the unexplained information (measured by the loss) through a model update.

GD IN ML AND PSEUDO-RESIDUALS

- By using the chain rule we see that the pseudo-residuals are input to the update direction

$$\begin{aligned}\nabla_{\theta} \mathcal{R}_{\text{emp}}(\theta) &= \sum_{i=1}^n \underbrace{\frac{\partial L(y^{(i)}, f)}{\partial f} \bigg|_{f=f(\mathbf{x}^{(i)} | \theta)}}_{=-\tilde{r}^{(i)}} \cdot \nabla_{\theta} f(\mathbf{x}^{(i)} | \theta) \\ &= - \sum_{i=1}^n \tilde{r}^{(i)} \cdot \nabla_{\theta} f(\mathbf{x}^{(i)} | \theta).\end{aligned}$$

- The unexplained information – the negative gradient – can be thought of as residuals, which is therefore also called pseudo-residuals.