


# Introduction to Machine Learning

## Nonlinear and Bayes



### Learning goals

● XXX

● XXX

# SUMMARY: REGULARIZED RISK MINIMIZATION

In  $\mathcal{R}_{\text{reg}}$  one has extreme flexibility to make appropriate choices

$$\mathcal{R}_{\text{reg}}(f) = \min_{f \in \mathcal{H}} \sum_{i=1}^n L\left(y^{(i)}, f\left(\mathbf{x}^{(i)}\right)\right) + \lambda \cdot J(f)$$

for a given ML problem:

- the **representation** of  $f$ , which determines how features can influence the predicted  $y$
- the **loss** function, which measures how errors should be treated
- the **regularization**  $J(f)$ , which encodes our inductive bias and preference for certain simpler models

By varying these choices one can construct a huge number of different ML models. Many ML models follow this construction principle or can be interpreted through the lens of regularized risk minimization.

# Regularization from a Bayesian Perspective

# REGULARIZED RISK MINIMIZATION VS. BAYES

We have already created a link between maximum likelihood estimation and empirical risk minimization.

Now we will generalize this for regularized risk minimization.

Assume we have a parameterized distribution  $p(\mathbf{x}|\theta)$  for our data and a prior  $p(\theta)$  over our parameter space, all in the Bayesian framework.

With Bayes theorem we know:

$$p(\theta|\mathbf{x}) = \frac{p(\mathbf{x}|\theta)p(\theta)}{p(\mathbf{x})} \propto p(\mathbf{x}|\theta)p(\theta)$$

# REGULARIZED RISK MINIMIZATION VS. BAYES


The maximum a posteriori (MAP) estimator of  $\theta$  is now the minimizer of

$$-\sum_{i=1}^n \log p(\mathbf{x}^{(i)} | \theta) - \log p(\theta).$$

Again, we identify the loss  $L(y, f(\mathbf{x} | \theta))$  with  $-\log(p(\mathbf{x}|\theta))$ . If  $p(\theta)$  is constant (i.e., we used a uniform, non-informative prior), we arrive at empirical risk minimization.

If not, we can identify  $J(\theta) \propto -\log(p(\theta))$ , i.e., the log-prior corresponds to the regularizer, and the additional control parameter  $\lambda$  corresponds to the relative strength of the prior in regularized risk minimization.

# REGULARIZED RISK MINIMIZATION VS. BAYES



- $L_2$  regularization corresponds to a zero-centered Gaussian prior,  $\theta_i \sim \mathcal{N}(0, \sigma^2)$ .
- $L_1$  regularization corresponds to a zero-centered Laplace prior,  $\theta_i \sim \text{Laplace}(0, b) = \frac{1}{2b} \exp(-\frac{|\theta_i|}{b})$ , where  $b$  is a scale parameter.
- In both cases, as regularization strength  $\lambda$  increases, the variance of the prior decreases, which in turn shrinks the parameters.